

```
In [2]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

In [4]: credit_card_data = pd.read_csv('creditcard.csv')
credit_card_data.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0

5 rows × 31 columns

```
In [6]: credit_card_data.sample()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
84047	60140.0	1.057529	-0.127356	0.900452	1.577774	-0.565966	0.298478	-0.332411	0.156855	0.846328	...	-0.064373	0.122969	-0.100207	0.115336	0.60709	-0.262013	0.067994	0.0285	40.0	0

1 rows × 31 columns

```
In [8]: credit_card_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype  
---  --
0   Time        284807 non-null  float64
1   V1          284807 non-null  float64
2   V2          284807 non-null  float64
3   V3          284807 non-null  float64
4   V4          284807 non-null  float64
5   V5          284807 non-null  float64
6   V6          284807 non-null  float64
7   V7          284807 non-null  float64
8   V8          284807 non-null  float64
9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
18  V18         284807 non-null  float64
19  V19         284807 non-null  float64
20  V20         284807 non-null  float64
21  V21         284807 non-null  float64
22  V22         284807 non-null  float64
23  V23         284807 non-null  float64
24  V24         284807 non-null  float64
25  V25         284807 non-null  float64
26  V26         284807 non-null  float64
27  V27         284807 non-null  float64
28  V28         284807 non-null  float64
29  Amount      284807 non-null  float64
30  Class       284807 non-null  int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

In [10]: credit_card_data.isnull().sum()
```

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
Amount	0
Class	0
dtype:	int64

```
In [12]: credit_card_data['Class'].value_counts()

Out[12]: Class
0      284315
1         492
Name: count, dtype: int64
```

```
In [14]: legit = credit_card_data[credit_card_data.Class==0]
fraud = credit_card_data[credit_card_data['Class']==1]
```

```
In [16]: fraud['Class']

Out[16]: 541      1
623      1
4920     1
6108     1
6329     1
      ..
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 492, dtype: int64
```

```
In [18]: legit.Amount.describe()

Out[18]: count      284315.000000
mean         88.291022
std          250.165092
min           0.000000
25%           5.650000
50%          22.000000
75%          77.050000
max         25691.160000
Name: Amount, dtype: float64
```

```
In [20]: fraud.Amount.describe()

Out[20]: count         492.000000
mean        122.211321
std         256.683288
min           0.000000
25%          1.000000
50%           9.250000
75%         105.890000
max         2125.870000
Name: Amount, dtype: float64
```

```
In [22]: credit_card_data.groupby('Class').mean()

Out[22]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount
Class																					
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637	-0.000987	0.004467	...	-0.000644	-0.001235	-0.000024	0.000070	0.000182	-0.000072	-0.000089	-0.000295	-0.000131	88.291022
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575	0.075667	122.211321

2 rows × 30 columns

```
In [24]: legit_sample = legit.sample(n=492)

In [26]: Df = pd.concat([legit_sample, fraud], axis=0)

In [28]: Df
```

[28]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
119522	75481.0	-0.410341	1.118510	0.488614	-0.037651	0.176083	-0.349123	0.449458	0.411255	-1.116376	...	0.178869	0.434948	-0.014137	0.069851	-0.272335	0.270254	-0.105427	-0.045919	10.62	0
202192	134210.0	2.077038	-0.014961	-2.100799	0.119556	0.582078	-1.121689	0.593713	-0.369255	-0.131720	...	0.101643	0.297303	-0.042122	-0.446975	0.242132	0.625305	-0.127526	-0.093778	21.74	0
257602	158249.0	-0.794127	-1.151332	-2.608456	0.315455	-1.217514	-0.291114	2.212935	-0.128182	-2.527802	...	0.912803	1.402599	1.276140	-0.432939	-0.719215	0.301387	-0.044643	0.199628	610.98	0
199451	132974.0	0.356572	-0.134729	-1.499157	-4.488220	0.254756	-0.832070	0.357918	-2.087665	0.458446	...	1.263518	-0.644885	-0.191495	0.021998	1.009419	-0.897916	0.311652	0.277715	59.99	0
191456	129264.0	-0.128263	-0.065892	-0.873523	-0.722465	1.502633	-0.410419	0.670331	-0.043873	0.072201	...	0.005265	0.084429	0.634069	0.083833	-1.968686	-0.104937	0.176572	0.367159	21.66	0
...
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...	0.778584	-0.319189	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00	1
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...	0.370612	0.028234	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76	1
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...	0.751826	0.834108	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89	1
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...	0.583276	-0.269209	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00	1
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...	-0.164350	-0.295135	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53	1

984 rows × 31 columns

```
In [30]: Df['Class'].value_counts()

Out[30]: Class
0      492
1      492
Name: count, dtype: int64
```

```
In [32]: Df.groupby('Class').mean()

Out[32]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount
Class																					
0	96664.479675	-0.032894	-0.059021	-0.027920	-0.030430	-0.076448	0.001536	0.104060	0.023887	0.038025	...	0.055565	-0.043044	-0.002856	-0.051717	-0.022704	-0.026194	0.000657	0.000499	0.005631	116.613394
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575	0.075667	122.211321

2 rows × 30 columns

```
In [34]: X = Df.drop(columns='Class', axis=1)
y = Df['Class']
```

```
In [36]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=2)
```

```
In [40]: model=LogisticRegression()
model.fit(X_train, y_train)

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(

Out[40]: LogisticRegression
LogisticRegression()
```

```
In [42]: X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, y_train)
print('Accuracy on Training data : ', training_data_accuracy)

Accuracy on Training data : 0.9504447268106735
```

```
In [44]: X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, y_test)
print('Accuracy score on Test Data : ', test_data_accuracy)

Accuracy score on Test Data : 0.9441624365482234
```

```
In [46]: Df.head()

Out[46]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
119522	75481.0	-0.410341	1.118510	0.488614	-0.037651	0.176083	-0.349123	0.449458	0.411255	-1.116376	...	0.178869	0.434948	-0.014137	0.069851	-0.272335	0.270254	-0.105427	-0.045919	10.62	0
202192	134210.0	2.077038	-0.014961	-2.100799	0.119556	0.582078	-1.121689	0.593713	-0.369255	-0.131720	...	0.101643	0.297303	-0.042122	-0.446975	0.242132	0.625305	-0.127526	-0.093778	21.74	0
257602	158249.0	-0.794127	-1.151332	-2.608456	0.315455	-1.217514	-0.291114	2.212935	-0.128182	-2.527802	...	0.912803	1.402599	1.276140	-0.432939	-0.719215	0.301387	-0.044643	0.199628	610.98	0
199451	132974.0	0.356572	-0.134729	-1.499157	-4.488220	0.254756	-0.832070	0.357918	-2.087665	0.458446	...	1.263518	-0.644885	-0.191495	0.021998	1.009419	-0.897916	0.311652	0.277715	59.99	0

191456	129264.0	-0.128263	-0.065892	-0.873523	-0.722465	1.502633	-0.410419	0.670331	-0.043873	0.072201	...	0.005265	0.084429	0.634069	0.083833	-1.968686	-0.104937	0.176572	0.367159	21.66	0
--------	----------	-----------	-----------	-----------	-----------	----------	-----------	----------	-----------	----------	-----	----------	----------	----------	----------	-----------	-----------	----------	----------	-------	---

5 rows × 31 columns